



.NET Blog

A first-hand look from the .NET engineering teams

Visual Studio

System.IO.Pipelines: High performance IO in .NET



July 9, 2018 by [David Fowler \(MSFT\)](#) // [17 Comments](#)



System.IO.Pipelines is a new library that is designed to make it easier to do high performance IO in .NET. It's a library targeting .NET Standard that works on all .NET implementations.

Pipelines was born from the work the .NET Core team did to make Kestrel one of the [fastest web servers in the industry](#). What started as an implementation detail inside of Kestrel progressed into a re-usable API that shipped in 2.1 as a first class BCL API (System.IO.Pipelines) available for all .NET developers.

What problem does it solve?

Correctly parsing data from a stream or socket is dominated by boilerplate code and has many corner cases, leading to complex code that is difficult to maintain.

Achieving high performance and being correct, while also dealing with this complexity is difficult. Pipelines aims to solve this complexity.

What extra complexity exists today?

Let's start with a simple problem. We want to write a TCP server that receives line-delimited messages (delimited by \n) from a client.

TCP Server with NetworkStream

DISCLAIMER: As with all performance sensitive work, each of the scenarios should be measured

within the context of your application. The overhead of the various techniques mentioned may not be necessary depending on the scale your networking applications need to handle.

The typical code you would write in .NET before pipelines looks something like this:

```
1  async Task ProcessLinesAsync(NetworkStream stream)
2  {
3      var buffer = new byte[1024];
4      await stream.ReadAsync(buffer, 0, buffer.Length);
5
6      // Process a single line from the buffer
7      ProcessLine(buffer);
8  }
```

sample1.cs hosted with ❤ by GitHub

[view raw](#)

This code might work when testing locally but it's has several errors:

- The entire message (end of line) may not have been received in a single call to `ReadAsync`.
- It's ignoring the result of `stream.ReadAsync()` which returns how much data was actually filled into the buffer.
- It doesn't handle the case where multiple lines come back in a single `ReadAsync` call.

These are some of the common pitfalls when reading streaming data. To account for this we need to make a few changes:

- We need to buffer the incoming data until we have found a new line.
- We need to parse *all* of the lines returned in the buffer

```
1  async Task ProcessLinesAsync(NetworkStream stream)
2  {
3      var buffer = new byte[1024];
4      var bytesBuffered = 0;
5      var bytesConsumed = 0;
6
7      while (true)
8      {
9          var bytesRead = await stream.ReadAsync(buffer, bytesBuffered, buffer.Length);
10         if (bytesRead == 0)
```

```

1      {
2          // EOF
3          break;
4      }
5      // Keep track of the amount of buffered bytes
6      bytesBuffered += bytesRead;
7
8      var linePosition = -1;
9
10     do
11     {
12         // Look for a EOL in the buffered data
13         linePosition = Array.IndexOf(buffer, (byte)'\n', bytesConsumed, bytesConsumed + bytesBuffered);
14
15         if (linePosition >= 0)
16         {
17             // Calculate the length of the line based on the offset
18             var lineLength = linePosition - bytesConsumed;
19
20             // Process the line
21             ProcessLine(buffer, bytesConsumed, lineLength);
22
23             // Move the bytesConsumed to skip past the line we consumed (including the line terminator)
24             bytesConsumed += lineLength + 1;
25         }
26     }
27     while (linePosition >= 0);
28 }
29

```

sample2.cs hosted with ❤ by GitHub

[view raw](#)

Once again, this might work in local testing but it's possible that the line is bigger than 1KiB (1024 bytes). We need to resize the input buffer until we have found a new line.

Also, we're allocating buffers on the heap as longer lines are processed. We can improve this by using the `ArrayPool<byte>` to avoid repeated buffer allocations as we parse longer lines from the client.

```

1  async Task ProcessLinesAsync(NetworkStream stream)
2  {
3      byte[] buffer = ArrayPool<byte>.Shared.Rent(1024);
4      var bytesBuffered = 0;
5      var bytesConsumed = 0;
6
7      while (true)
8      {
9          // Calculate the amount of bytes remaining in the buffer
10         var bytesRemaining = buffer.Length - bytesBuffered;
11
12         if (bytesRemaining == 0)
13         {
14             // Double the buffer size and copy the previously buffered data into
15             var newBuffer = ArrayPool<byte>.Shared.Rent(buffer.Length * 2);
16             Buffer.BlockCopy(buffer, 0, newBuffer, 0, buffer.Length);
17             // Return the old buffer to the pool
18             ArrayPool<byte>.Shared.Return(buffer);
19             buffer = newBuffer;
20             bytesRemaining = buffer.Length - bytesBuffered;
21         }
22
23         var bytesRead = await stream.ReadAsync(buffer, bytesBuffered, bytesRemaining);
24         if (bytesRead == 0)
25         {
26             // EOF
27             break;
28         }
29
30         // Keep track of the amount of buffered bytes
31         bytesBuffered += bytesRead;
32
33         do
34         {
35             // Look for a EOL in the buffered data
36             linePosition = Array.IndexOf(buffer, (byte)'\n', bytesConsumed, bytesRemaining);
37
38             if (linePosition >= 0)

```

```

39         {
40             // Calculate the length of the line based on the offset
41             var lineLength = linePosition - bytesConsumed;
42
43             // Process the line
44             ProcessLine(buffer, bytesConsumed, lineLength);
45
46             // Move the bytesConsumed to skip past the line we consumed (in
47             bytesConsumed += lineLength + 1;
48         }
49     }
50     while (linePosition >= 0);
51 }
52 }

```

sample3.cs hosted with ❤ by GitHub

[view raw](#)

This code works but now we're re-sizing the buffer which results in more buffer copies. It also uses more memory as the logic doesn't shrink the buffer after lines are processed. To avoid this, we can store a list of buffers instead of resizing each time we cross the 1KiB buffer size.

Also, we don't grow the 1KiB buffer until it's completely empty. This means we can end up passing smaller and smaller buffers to `ReadAsync` which will result in more calls into the operating system.

To mitigate this, we'll allocate a new buffer when there's less than 512 bytes remaining in the existing buffer:

```

1  public class BufferSegment
2  {
3      public byte[] Buffer { get; set; }
4      public int Count { get; set; }
5
6      public int Remaining => Buffer.Length - Count;
7  }
8
9  async Task ProcessLinesAsync(NetworkStream stream)
10 {
11     const int minimumBufferSize = 512;

```

```

12
13     var segments = new List<BufferSegment>();
14     var bytesConsumed = 0;
15     var bytesConsumedBufferIndex = 0;
16     var segment = new BufferSegment { Buffer = ArrayPool<byte>.Shared.Rent(1024);
17
18     segments.Add(segment);
19
20     while (true)
21     {
22         // Calculate the amount of bytes remaining in the buffer
23         if (segment.Remaining < minimumBufferSize)
24         {
25             // Allocate a new segment
26             segment = new BufferSegment { Buffer = ArrayPool<byte>.Shared.Rent(1024);
27             segments.Add(segment);
28         }
29
30         var bytesRead = await stream.ReadAsync(segment.Buffer, segment.Count, CancellationToken.None);
31         if (bytesRead == 0)
32         {
33             break;
34         }
35
36         // Keep track of the amount of buffered bytes
37         segment.Count += bytesRead;
38
39         while (true)
40         {
41             // Look for a EOL in the list of segments
42             var (segmentIndex, segmentOffset) = IndexOf(segments, (byte)'\n', bytesConsumedBufferIndex);
43
44             if (segmentIndex >= 0)
45             {
46                 // Process the line
47                 ProcessLine(segments, segmentIndex, segmentOffset);
48
49                 bytesConsumedBufferIndex = segmentOffset;

```

```

50         bytesConsumed = segmentOffset + 1;
51     }
52     else
53     {
54         break;
55     }
56 }
57
58 // Drop fully consumed segments from the list so we don't look at them
59 for (var i = bytesConsumedBufferIndex; i >= 0; --i)
60 {
61     var consumedSegment = segments[i];
62     // Return all segments unless this is the current segment
63     if (consumedSegment != segment)
64     {
65         ArrayPool<byte>.Shared.Return(consumedSegment.Buffer);
66         segments.RemoveAt(i);
67     }
68 }
69 }
70 }
71
72 (int segmentIndex, int segmentOffset) IndexOf(List<BufferSegment> segments, byte
73 {
74     var first = true;
75     for (var i = startBufferIndex; i < segments.Count; ++i)
76     {
77         var segment = segments[i];
78         // Start from the correct offset
79         var offset = first ? startSegmentOffset : 0;
80         var index = Array.IndexOf(segment.Buffer, value, offset, segment.Count
81
82         if (index >= 0)
83         {
84             // Return the buffer index and the index within that segment where
85             return (i, index);
86         }
87

```

```
88         first = false;
89     }
90     return (-1, -1);
91 }
```

sample4.cs hosted with ❤ by GitHub

[view raw](#)

This code just got *much* more complicated. We're keeping track of the filled up buffers as we're looking for the delimiter. To do this, we're using a `List<BufferSegment>` here to represent the buffered data while looking for the new line delimiter. As a result, `ProcessLine` and `IndexOf` now accept a `List<BufferSegment>` instead of a `byte[]`, `offset` and `count`. Our parsing logic needs to now handle one or more buffer segments.

Our server now handles partial messages, and it uses pooled memory to reduce overall memory consumption but there are still a couple more changes we need to make:

1. The `byte[]` we're using from the `ArrayPool<byte>` are just regular managed arrays. This means whenever we do a `ReadAsync` or `WriteAsync`, those buffers get pinned for the lifetime of the asynchronous operation (in order to interop with the native IO APIs on the operating system). This has performance implications on the garbage collector since pinned memory cannot be moved which can lead to heap fragmentation. Depending on how long the async operations are pending, the pool implementation may need to change.
2. The throughput can be optimized by decoupling the reading and processing logic. This creates a batching effect that lets the parsing logic consume larger chunks of buffers, instead of reading more data only after parsing a single line. This introduces some additional complexity:
 - We need two loops that run independently of each other. One that reads from the `Socket` and one that parses the buffers.
 - We need a way to signal the parsing logic when data becomes available.
 - We need to decide what happens if the loop reading from the `Socket` is "too fast". We need a way to throttle the reading loop if the parsing logic can't keep up. This is commonly referred to as "flow control" or "back pressure".
 - We need to make sure things are thread safe. We're now sharing a set of buffers between the reading loop and the parsing loop and those run independently on different threads.
 - The memory management logic is now spread across two different pieces of code, the code that rents from the buffer pool is reading from the socket and the code that returns from the buffer pool is the parsing logic.
 - We need to be extremely careful with how we return buffers after the parsing logic is

done with them. If we're not careful, it's possible that we return a buffer that's still being written to by the `Socket` reading logic.

The complexity has gone through the roof (and we haven't even covered all of the cases). High performance networking usually means writing very complex code in order to eke out more performance from the system.

The goal of `System.IO.Pipelines` is to make writing this type of code easier.

TCP server with `System.IO.Pipelines`

Let's take a look at what this example looks like with `System.IO.Pipelines`:

```
1  async Task ProcessLinesAsync(Socket socket)
2  {
3      var pipe = new Pipe();
4      Task writing = FillPipeAsync(socket, pipe.Writer);
5      Task reading = ReadPipeAsync(pipe.Reader);
6
7      return Task.WhenAll(reading, writing);
8  }
9
10 async Task FillPipeAsync(Socket socket, PipeWriter writer)
11 {
12     const int minimumBufferSize = 512;
13
14     while (true)
15     {
16         // Allocate at least 512 bytes from the PipeWriter
17         Memory<byte> memory = writer.GetMemory(minimumBufferSize);
18         try
19         {
20             int bytesRead = await socket.ReceiveAsync(memory, SocketFlags.None);
21             if (bytesRead == 0)
22             {
23                 break;
24             }
25             // Tell the PipeWriter how much was read from the Socket
26             writer.Advance(bytesRead);
```

```
27     }
28     catch (Exception ex)
29     {
30         LogError(ex);
31         break;
32     }
33
34     // Make the data available to the PipeReader
35     FlushResult result = await writer.FlushAsync();
36
37     if (result.IsCompleted)
38     {
39         break;
40     }
41 }
42
43 // Tell the PipeReader that there's no more data coming
44 writer.Complete();
45 }
46
47 async Task ReadPipeAsync(PipeReader reader)
48 {
49     while (true)
50     {
51         ReadResult result = await reader.ReadAsync();
52
53         ReadOnlySequence<byte> buffer = result.Buffer;
54         SequencePosition? position = null;
55
56         do
57         {
58             // Look for a EOL in the buffer
59             position = buffer.PositionOf((byte)'\n');
60
61             if (position != null)
62             {
63                 // Process the line
64                 ProcessLine(buffer.Slice(0, position.Value));
```

```

65
66         // Skip the line + the \n character (basically position)
67         buffer = buffer.Slice(buffer.GetPosition(1, position.Value));
68     }
69 }
70 while (position != null);
71
72 // Tell the PipeReader how much of the buffer we have consumed
73 reader.AdvanceTo(buffer.Start, buffer.End);
74
75 // Stop reading if there's no more data coming
76 if (result.IsCompleted)
77 {
78     break;
79 }
80 }
81
82 // Mark the PipeReader as complete
83 reader.Complete();
84 }

```

sample5.cs hosted with ❤ by GitHub

[view raw](#)

The pipelines version of our line reader has 2 loops:

- `FillPipeAsync` reads from the `Socket` and writes into the `PipeWriter`.
- `ReadPipeAsync` reads from the `PipeReader` and parses incoming lines.

Unlike the original examples, there are no explicit buffers allocated anywhere. This is one of pipelines' core features. All buffer management is delegated to the `PipeReader`/`PipeWriter` implementations.

This makes it easier for consuming code to focus solely on the business logic instead of complex buffer management.

In the first loop, we first call `PipeWriter.GetMemory(int)` to get some memory from the underlying writer; then we call `PipeWriter.Advance(int)` to tell the `PipeWriter` how much data we actually wrote to the buffer. We then call `PipeWriter.FlushAsync()` to make the data available to the `PipeReader`.

In the second loop, we're consuming the buffers written by the `PipeWriter` which ultimately comes from the `Socket`. When the call to `PipeReader.ReadAsync()` returns, we get a `ReadResult` which contains 2 important pieces of information, the data that was read in the form of `ReadOnlySequence<byte>` and a bool `IsCompleted` that lets the reader know if the writer is done writing (EOF). After finding the end of line (EOL) delimiter and parsing the line, we slice the buffer to skip what we've already processed and then we call `PipeReader.AdvanceTo` to tell the `PipeReader` how much data we have consumed.

At the end of each of the loops, we complete both the reader and the writer. This lets the underlying `Pipe` release all of the memory it allocated.

System.IO.Pipelines

Partial Reads

Besides handling the memory management, the other core pipelines feature is the ability to peek at data in the `Pipe` without actually consuming it.

`PipeReader` has two core APIs `ReadAsync` and `AdvanceTo`. `ReadAsync` gets the data in the `Pipe`, `AdvanceTo` tells the `PipeReader` that these buffers are no longer required by the reader so they can be discarded (for example returned to the underlying buffer pool).

Here's an example of an http parser that reads partial data buffers data in the `Pipe` until a valid start line is received.

The reader can look at data without consuming it.

```
Pipe.Reader.ReadAsync()
```

GE

```
HttpParser.TryParse( GE )
```

false

```
Pipe.Reader.AdvanceTo( GE ).Start()
```

```
Pipe.Reader.ReadAsync()
```

GET /\r\n

```
HttpParser.TryParse( GET /\r\n )
```

true

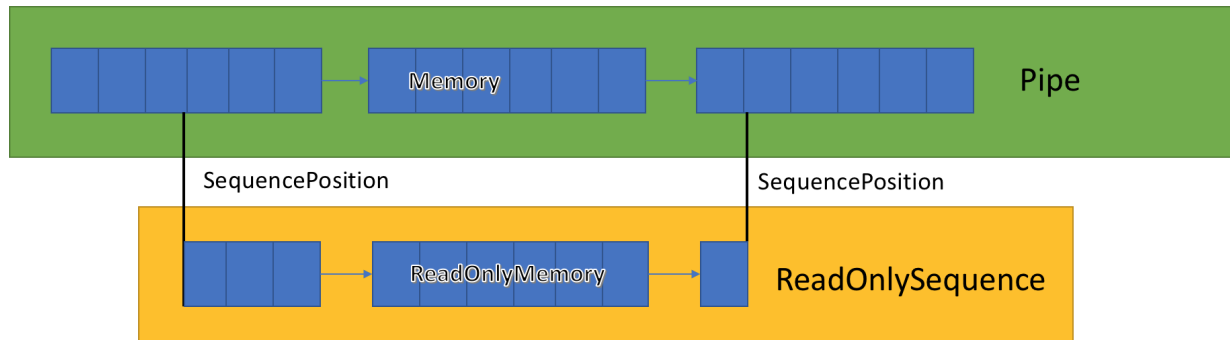
```
Pipe.Reader.AdvanceTo( GET /\r\n ).End()
```

```
Pipe.Reader.ReadAsync()
```

Host: localhost ...

ReadOnlySequence<T>

The `Pipe` implementation stores a linked list of buffers that get passed between the `PipeWriter` and `PipeReader`. `PipeReader.ReadAsync` exposes a `ReadOnlySequence<T>` which is a new BCL type that represents a view over one or more segments of `ReadOnlyMemory<T>`, similar to `Span<T>` and `Memory<T>` which provide a view over arrays and strings.



The `Pipe` internally maintains pointers to where the reader and writer are in the overall set of allocated data and updates them as data is written or read. The `SequencePosition` represents a single point in the linked list of buffers and can be used to efficiently slice the `ReadOnlySequence<T>`.

Since the `ReadOnlySequence<T>` can support one or more segments, it's typical for high performance processing logic to split fast and slow paths based on single or multiple segments.

For example, here's a routine that converts an ASCII `ReadOnlySequence<byte>` into a `string`:

```
1  string GetAsciiString(ReadOnlySequence<byte> buffer)
2  {
3      if (buffer.IsSingleSegment)
4      {
5          return Encoding.ASCII.GetString(buffer.First.Span);
6      }
7
8      return string.Create((int)buffer.Length, buffer, (span, sequence) =>
9      {
10         foreach (var segment in sequence)
11         {
```

```

12         Encoding.ASCII.GetChars(segment.Span, span);
13
14         span = span.Slice(segment.Length);
15     }
16 });
17 }

```


sample6.cs hosted with ❤ by GitHub

[view raw](#)

Back pressure and flow control

In a perfect world, reading & parsing work as a team: the reading thread consumes the data from the network and puts it in buffers while the parsing thread is responsible for constructing the appropriate data structures. Normally, parsing will take more time than just copying blocks of data from the network. As a result, the reading thread can easily overwhelm the parsing thread. The result is that the reading thread will have to either slow down or allocate more memory to store the data for the parsing thread. For optimal performance, there is a balance between frequent pauses and allocating more memory.

To solve this problem, the pipe has two settings to control the flow of data, the `PauseWriterThreshold` and the `ResumeWriterThreshold`. The `PauseWriterThreshold` determines how much data should be buffered before calls to `PipeWriter.FlushAsync` pauses. The `ResumeWriterThreshold` controls how much the reader has to consume before writing can resume.



`PipeWriter.FlushAsync` “blocks” when the amount of data in the Pipe crosses `PauseWriterThreshold` and “unblocks” when it becomes lower than `ResumeWriterThreshold`. Two values are used to prevent thrashing around the limit.

Scheduling IO

Usually when using `async/await`, continuations are called on either on thread pool threads or on the

current `SynchronizationContext`.

When doing IO it's very important to have fine-grained control over where that IO is performed so that one can take advantage of CPU caches more effectively, which is critical for high-performance applications like web servers. `Pipelines` exposes a `PipeScheduler` that determines where asynchronous callbacks run. This gives the caller fine-grained control over exactly what threads are used for IO.

An example of this in practice is in the Kestrel Libuv transport where IO callbacks run on dedicated event loop threads.

Other benefits of the `PipeReader` pattern:

- Some underlying systems support a “bufferless wait”, that is, a buffer never needs to be allocated until there's actually data available in the underlying system. For example on Linux with `epoll`, it's possible to wait until data is ready before actually supplying a buffer to do the read. This avoids the problem where having a large number of threads waiting for data doesn't immediately require reserving a huge amount of memory.
- The default `Pipe` makes it easy to write unit tests against networking code because the parsing logic is separated from the networking code so unit tests only run the parsing logic against in-memory buffers rather than consuming directly from the network. It also makes it easy to test those hard to test patterns where partial data is sent. ASP.NET Core uses this to test various aspects of the Kestrel's http parser.
- Systems that allow exposing the underlying OS buffers (like the Registered IO APIs on Windows) to user code are a natural fit for pipelines since buffers are always provided by the `PipeReader` implementation.

Other Related types

As part of making `System.IO.Pipelines`, we also added a number of new primitive BCL types:

- [MemoryPool<T>](#), [IMemoryOwner<T>](#), [MemoryManager<T>](#) – .NET Core 1.0 added [ArrayPool<T>](#) and in .NET Core 2.1 we now have a more general abstraction for a pool that works over any `Memory<T>`. This provides an extensibility point that lets you plug in more advanced allocation strategies as well as control how buffers are managed (for e.g. provide pre-pinned buffers instead of purely managed arrays).
- [IBufferWriter<T>](#) – Represents a sink for writing synchronous buffered data. (`PipeWriter` implements this)
- [IValueTaskSource](#) – [ValueTask<T>](#) has existed since .NET Core 1.1 but has gained some super powers in .NET Core 2.1 to allow allocation-free awaitable async operations. See

<https://github.com/dotnet/corefx/issues/27445> for more details.

How do I use Pipelines?

The APIs exist in the `System.IO.Pipelines` nuget package.

Here's an example of a .NET Core 2.1 server application that uses pipelines to handle line based messages (our example above) <https://github.com/davidfowl/TcpEcho>. It should run with `dotnet run` (or by running it in Visual Studio). It listens to a socket on port 8087 and writes out received messages to the console. You can use a client like netcat or putty to make a connection to 8087 and send line based messages to see it working.

Today Pipelines powers Kestrel and SignalR, and we hope to see it at the center of many networking libraries and components from the .NET community.



.NET Application Architecture

[Microservices with Docker Containers](#)

[Web apps with ASP.NET Core](#)

[Mobile apps with Xamarin.Forms](#)

[Modernizing existing .NET apps to the cloud](#)

☐ Search this blog ☐ Search all blogs

Tags

.NET [.NET Architecture](#) **.NET Core** [.net framework](#) [.NET Update](#)
announcement [asp.net](#) [ASP.NET Core](#) [async](#) [azure](#) [bcl](#) [c#](#) [clr](#) [codegen](#)
community [compiler](#) [conferences](#) [diagnostics](#) [dotnetnative](#) [Entity](#)

[Framework](#) [Entity Framework Core](#) [F# fundamentals](#) [httpclient](#) [jit](#) [nuget](#) [On .NET](#)
[open source](#) [performance](#) [portable](#) [class libraries](#) [releases](#) [roslyn](#) [ryujit](#)
[Security Update Servicing](#) [vb](#) [Visual F#](#) [visual studio](#) **Week in .NET** [windows](#)
[store](#) [Windows Update](#) [wpf](#) [WSUS](#) [WU](#) [xamarin](#)

Archives

[July 2018](#) (4)
[June 2018](#) (8)
[May 2018](#) (13)
[April 2018](#) (8)
[March 2018](#) (3)
[February 2018](#) (9)
[January 2018](#) (5)
[All of 2018](#) (50)
[All of 2017](#) (119)
[All of 2016](#) (119)
[All of 2015](#) (40)
[All of 2014](#) (56)
[All of 2013](#) (40)
[All of 2012](#) (27)
[All of 2011](#) (6)
[All of 2010](#) (12)
[All of 2009](#) (52)

Join the conversation

Add Comment



Mark Eastwood

2 days ago



Thanks for the great blog David!

I've been writing some TCP code based on SocketAsyncEventArgs (<https://docs.microsoft.com/en-us/dotnet/api/system.net.sockets.socketasyncEventArgs?view=netcore-2.1>). Will I see much of a performance/memory allocation improvement using Pipelines?





David Fowler

1 day ago



It definitely *can**. It'll probably help you write more correct code first off.

It'll also force you into patterns that can improve your memory allocations and throughput (as described in the article). I'd definitely ask you to try the API out and see what you end up with.



Alvaro Rivoir

2 days ago



How should we consume network streams and files using Pipelines?



Vince

2 days ago



Thank you. BTW, "delimeter" should be "delimiter"



David Fowler

1 day ago



Thanks!



PetSerAI

2 days ago



Should it be ``bytesConsumedBufferIndex = segmentIndex;`` instead of ``bytesConsumedBufferIndex = segmentOffset;``?



Łukasz Pyrzyk

1 day ago



Great post!



Derekk

1 day ago



Excellent!

Marcel Veldhuizen

1 day ago



The given code assumes ASCII encoding when searching for the newline. A short note about that might have been good, before people start blindly copying this code 😊

I understand actually fixing it is outside the scope of the example, unless there is a clever way of dealing with (for example) UTF-8 that I'm not aware of...



Marc Stromko

23 hours ago



The newline sequence is the same in UTF-8, all 7 bit ASCII characters are the same in the UTF formats. For UTF-8 support beyond ASCII, change the Encoding.ASCII to Encoding.UTF8, and verify that the read buffer is not ending mid-character.



paradyne at work

22 hours ago



There is a clever way.. Or rather UTF-8 is clever. It never uses values of 0-127 inside multibyte characters so you are safe to just search through to a newline and then convert the bytes up to that point into a string using UTF8.GetString().

See: <https://en.wikipedia.org/wiki/UTF-8>



Kevin Bryant

24 hours ago



Thanks for this! Can we please have a post like this for `system.threading.channels`? The readme from the `corefxlabs` repo seems to be the only existing documentation, and it appears to be out of date. The comments in the code are great, but aside from the hints in the early readme there doesn't seem to be much describing how it's all intended to fit together. Even a pointer to somewhere it's being used in another project – haven't had much luck finding usage so far.



Juliano Goncalves

21 hours ago



Really nice article!

I wonder what it would take to abstract this even more maybe using `IObservable`` and Rx though. I always found Rx's way of representing this kind of push operation to be incredibly streamlined and simple to comprehend.

Even though the new API is drastically simpler than the old version, I still find it somewhat complex and hard to read. Maybe if another abstraction level is introduced it would be possible to make this a little bit more intuitive.



Yandy Zaldivar

19 hours ago



I've been using a custom TCP binary protocol not delimited by lines or any other character(s), just binary serialized objects with length prefixed. Is this library able to handle this kind of behavior or only line-delimited messages?



Rafael Teixeira

15 hours ago



Definitively yes. You just keep reading until the length you need is completed. You can even get smart with asking buffers from the pipe that match the expected lengths.



Eugeniy

18 hours ago



Don't catch catch(exception) without of rethrowing



Steve S

12 hours ago



Great Post, Fantastic Library

© 2018 Microsoft Corporation.

[Terms of Use](#) | [Trademarks](#) | [Privacy & Cookies](#)

Microsoft